

DeepSeek V4—almost on the frontier, a fraction of the price

Chinese AI lab DeepSeek's last model release was V3.2 (and V3.2 Speciale) [last December](#). They just dropped the first of their hotly anticipated V4 series in the shape of two preview models, [DeepSeek-V4-Pro](#) and [DeepSeek-V4-Flash](#).

Both models are 1 million token context Mixture of Experts. Pro is 1.6T total parameters, 49B active. Flash is 284B total, 13B active. They're using the standard MIT license.

I think this makes DeepSeek-V4-Pro the new largest open weights model. It's larger than Kimi K2.6 (1.1T) and GLM-5.1 (754B) and more than twice the size of DeepSeek V3.2 (685B).


Pro is 865GB on Hugging Face, Flash is 160GB. I'm hoping that a lightly quantized Flash will run on my 128GB M5 MacBook Pro. It's *possible* the Pro model may run on it if I can stream just the necessary active experts from disk.

For the moment I tried the models out via [OpenRouter](#), using [llm-openrouter](#):

Bash

```
1 llm install llm-openrouter
2 llm openrouter refresh
3 llm -m openrouter/deepseek/deepseek-v4-pro 'Generate an SVG of a pelican riding a bicycle'
```

Here's the pelican [for DeepSeek-V4-Flash](#):



And for DeepSeek-V4-Pro:



For comparison, take a look at the pelicans I got from [DeepSeek V3.2 in December](#), [V3.1 in August](#), and [V3-0324 in March 2025](#).

So the pelicans are pretty good, but what's really notable here is the *cost*. DeepSeek V4 is a very, very inexpensive model.

This is [DeepSeek's pricing page](#). They're charging \$0.14/million tokens input and \$0.28/million tokens output for Flash, and \$1.74/million input and \$3.48/million output for Pro.

Here's a comparison table with the frontier models from Gemini, OpenAI and Anthropic:

Model	Input (\$/M)	Output (\$/M)
DeepSeek V4 Flash	\$0.14	\$0.28
GPT-5.4 Nano	\$0.20	\$1.25
Gemini 3.1 Flash-Lite	\$0.25	\$1.50
Gemini 3 Flash Preview	\$0.50	\$3
GPT-5.4 Mini	\$0.75	\$4.50
Claude Haiku 4.5	\$1	\$5
DeepSeek V4 Pro	\$1.74	\$3.48
Gemini 3.1 Pro	\$2	\$12
GPT-5.4	\$2.50	\$15
Claude Sonnet 4.6	\$3	\$15
Claude Opus 4.7	\$5	\$25
GPT-5.5	\$5	\$30

DeepSeek-V4-Flash is the cheapest of the small models, beating even OpenAI’ s GPT-5.4 Nano. DeepSeek-V4-Pro is the cheapest of the larger frontier models.

This note from [the DeepSeek paper](#) helps explain why they can price these models so low —they’ ve focused a great deal on efficiency with this release, especially for longer context prompts:

In the scenario of 1M-token context, even DeepSeek-V4-Pro, which has a larger number of activated parameters, attains only 27% of the single-token FLOPs (measured in equivalent FP8 FLOPs) and 10% of the KV cache size relative to DeepSeek-V3.2. Furthermore, DeepSeek-V4-Flash, with its smaller number of activated parameters, pushes efficiency even further: in the 1M-token context setting, it achieves only 10% of the single-token FLOPs and 7% of the KV cache size compared with DeepSeek-V3.2.

DeepSeek’ s self-reported benchmarks [in their paper](#) show their Pro model competitive with those other frontier models, albeit with this note:

Through the expansion of reasoning tokens, DeepSeek-V4-Pro-Max demonstrates superior performance relative to GPT-5.2 and Gemini-3.0-Pro on standard reasoning benchmarks. Nevertheless, its performance falls marginally short of GPT-5.4 and

Gemini-3.1-Pro, suggesting a developmental trajectory that trails state-of-the-art frontier models by approximately 3 to 6 months.

I'm keeping an eye on huggingface.co/unsloth/models as I expect the Unsloth team will have a set of quantized versions out pretty soon. It's going to be very interesting to see how well that Flash model runs on my own machine.

Posted [24th April 2026](#) at 6:01 am • Follow me on [Mastodon](#), [Bluesky](#), [Twitter](#) or [subscribe to my newsletter](#)

More recent articles

- [Extract PDF text in your browser with LiteParse for the web](#) - 23rd April 2026
- [A pelican for GPT-5.5 via the semi-official Codex backdoor API](#) - 23rd April 2026